

CPU Scheduling

CHAPTER 5



CPU scheduling is the basis of multiprogrammed operating systems. By switching the CPU among processes, the operating system can make the computer more productive. In this chapter, we introduce basic CPU-scheduling concepts and present several CPU-scheduling algorithms, including real-time systems. We also consider the problem of selecting an algorithm for a particular system.

In Chapter 4, we introduced threads to the process model. On modern operating systems it is kernel-level threads—not processes—that are in fact being scheduled by the operating system. However, the terms "process scheduling" and "thread scheduling" are often used interchangeably. In this chapter, we use *process scheduling* when discussing general scheduling concepts and *thread scheduling* to refer to thread-specific ideas.

Similarly, in Chapter 1 we describe how a *core* is the basic computational unit of a CPU, and that a process executes on a CPU's core. However, in many instances in this chapter, when we use the general terminology of scheduling a process to "run on a CPU", we are implying that the process is running on a CPU's core.

CHAPTER OBJECTIVES

- Describe various CPU scheduling algorithms.
- Assess CPU scheduling algorithms based on scheduling criteria.
- Explain the issues related to multiprocessor and multicore scheduling.
- Describe various real-time scheduling algorithms.
- Describe the scheduling algorithms used in the Windows, Linux, and Solaris operating systems.
- Apply modeling and simulations to evaluate CPU scheduling algorithms.
- Design a program that implements several different CPU scheduling algorithms.

5.1 Basic Concepts

In a system with a single CPU core, only one process can run at a time. Others must wait until the CPU's core is free and can be rescheduled. The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization. The idea is relatively simple. A process is executed until it must wait, typically for the completion of some I/O request. In a simple computer system, the CPU then just sits idle. All this waiting time is wasted; no useful work is accomplished. With multiprogramming, we try to use this time productively. Several processes are kept in memory at one time. When one process has to wait, the operating system takes the CPU away from that process and gives the CPU to another process. This pattern continues. Every time one process has to wait, another process can take over use of the CPU. On a multicore system, this concept of keeping the CPU busy is extended to all processing cores on the system.

Scheduling of this kind is a fundamental operating-system function. Almost all computer resources are scheduled before use. The CPU is, of course, one of the primary computer resources. Thus, its scheduling is central to operating-system design.




Figure 5.1 Alternating sequence of CPU and I/O bursts.

5.1.1 CPU-I/O Burst Cycle

The success of CPU scheduling depends on an observed property of processes: process execution consists of a **cycle** of CPU execution and I/O wait. Processes alternate between these two states. Process execution begins with a **CPU burst**. That is followed by an **I/O burst**, which is followed by another CPU burst, then another I/O burst, and so on. Eventually, the final CPU burst ends with a system request to terminate execution (Figure 5.1).

The durations of CPU bursts have been measured extensively. Although they vary greatly from process to process and from computer to computer, they tend to have a frequency curve similar to that shown in Figure 5.2. The curve is generally characterized as exponential or hyperexponential, with a large number of short CPU bursts and a small number of long CPU bursts. An I/O-bound program typically has many short CPU bursts. A CPU-bound program might have a few long CPU bursts. This distribution can be important when implementing a CPU-scheduling algorithm.

5.1.2 CPU Scheduler

Whenever the CPU becomes idle, the operating system must select one of the processes in the ready queue to be executed. The selection process is carried out by the **CPU scheduler**, which selects a process from the processes in memory that are ready to execute and allocates the CPU to that process.

Note that the ready queue is not necessarily a first-in, first-out (FIFO) queue. As we shall see when we consider the various scheduling algorithms, a ready queue can be implemented as a FIFO queue, a priority queue, a tree, or simply an unordered linked list. Conceptually, however, all the processes in the ready queue are lined up waiting for a chance to run on the CPU. The records in the queues are generally process control blocks (PCBs) of the processes.




Figure 5.2 Histogram of CPU-burst durations.

5.1.3 Preemptive and Nonpreemptive Scheduling

CPU-scheduling decisions may take place under the following four circumstances:

1. When a process switches from the running state to the waiting state (for example, as the result of an I/O request or an invocation of `wait()` for the termination of a child process)
2. When a process switches from the running state to the ready state (for example, when an interrupt occurs)
3. When a process switches from the waiting state to the ready state (for example, at completion of I/O)
4. When a process terminates

For situations 1 and 4, there is no choice in terms of scheduling. A new process (if one exists in the ready queue) must be selected for execution. There is a choice, however, for situations 2 and 3.

When scheduling takes place only under circumstances 1 and 4, we say that the scheduling scheme is **nonpreemptive** or **cooperative**. Otherwise, it is **preemptive**. Under nonpreemptive scheduling, once the CPU has been allocated to a process, the process keeps the CPU until it releases it either by terminating or by switching to the waiting state. Virtually all modern operating systems including Windows, macOS, Linux, and UNIX use preemptive scheduling algorithms.

Unfortunately, preemptive scheduling can result in race conditions when data are shared among several processes. Consider the case of two processes that share data. While one process is updating the data, it is preempted so that the second process can run. The second process then tries to read the data, which are in an inconsistent state. This issue will be explored in detail in Chapter 6.

Preemption also affects the design of the operating-system kernel. During the processing of a system call, the kernel may be busy with an activity on behalf of a process. Such activities may involve changing important kernel data (for instance, I/O queues). What happens if the process is preempted in the middle of these changes and the kernel (or the device driver) needs to read or modify the same structure? Chaos ensues. As will be discussed in Section 6.2, operating-system kernels can be designed as either nonpreemptive or preemptive. A nonpreemptive kernel will wait for a system call to complete or for a process to block while waiting for I/O to complete to take place before doing a context switch. This scheme ensures that the kernel structure is simple, since the kernel will not preempt a process while the kernel data structures are in an inconsistent state. Unfortunately, this kernel-execution model is a poor one for supporting real-time computing, where tasks must complete execution within a given time frame. In Section 5.6, we explore scheduling demands of real-time systems. A preemptive kernel requires mechanisms such as mutex locks to prevent race conditions when accessing shared kernel data structures. Most modern operating systems are now fully preemptive when running in kernel mode.




Figure 5.3 The role of the dispatcher.

Because interrupts can, by definition, occur at any time, and because they cannot always be ignored by the kernel, the sections of code affected by interrupts must be guarded from simultaneous use. The operating system needs to accept interrupts at almost all times. Otherwise, input might be lost or output overwritten. So that these sections of code are not accessed concurrently by several processes, they disable interrupts at entry and reenables interrupts at exit. It is important to note that sections of code that disable interrupts do not occur very often and typically contain few instructions.

5.1.4 Dispatcher

Another component involved in the CPU-scheduling function is the **dispatcher**. The dispatcher is the module that gives control of the CPU's core to the process selected by the CPU scheduler. This function involves the following:

- Switching context from one process to another
- Switching to user mode
- Jumping to the proper location in the user program to resume that program

The dispatcher should be as fast as possible, since it is invoked during every context switch. The time it takes for the dispatcher to stop one process and start another running is known as the **dispatch latency** and is illustrated in Figure 5.3.

An interesting question to consider is, how often do context switches occur? On a system-wide level, the number of context switches can be obtained by using the `vmstat` command that is available on Linux systems. Below is the output (which has been trimmed) from the command

```
vmstat 1 3
```

This command provides 3 lines of output over a 1-second delay:

```
-----cpu-----
24
225
339
```

The first line gives the average number of context switches over 1 second since the system booted, and the next two lines give the number of context switches over two 1-second intervals. Since this machine booted, it has averaged 24 context switches per second. And in the past second, 225 context switches were made, with 339 context switches in the second prior to that.

We can also use the `/proc` file system to determine the number of context switches for a given process. For example, the contents of the file `/proc/2166/status` will list various statistics for the process with `pid = 2166`. The command

```
cat /proc/2166/status
```

provides the following trimmed output:

```
voluntary_ctxt_switches      150
nonvoluntary_ctxt_switches   8
```

This output shows the number of context switches over the lifetime of the process. Notice the distinction between *voluntary* and *nonvoluntary* context switches. A voluntary context switch occurs when a process has given up control of the CPU because it requires a resource that is currently unavailable (such as blocking for I/O.) A nonvoluntary context switch occurs when the CPU has been taken away from a process, such as when its time slice has expired or it has been preempted by a higher-priority process.

5.2 Scheduling Criteria

Different CPU-scheduling algorithms have different properties, and the choice of a particular algorithm may favor one class of processes over another. In choosing which algorithm to use in a particular situation, we must consider the properties of the various algorithms.

Many criteria have been suggested for comparing CPU-scheduling algorithms. Which characteristics are used for comparison can make a substantial difference in which algorithm is judged to be best. The criteria include the following:

- **CPU utilization.** We want to keep the CPU as busy as possible. Conceptually, CPU utilization can range from 0 to 100 percent. In a real system, it should range from 40 percent (for a lightly loaded system) to 90 percent (for a heavily loaded system). (CPU utilization can be obtained by using the `top` command on Linux, macOS, and UNIX systems.)
- **Throughput.** If the CPU is busy executing processes, then work is being done. One measure of work is the number of processes that are completed

per time unit, called **throughput**. For long processes, this rate may be one process over several seconds; for short transactions, it may be tens of processes per second.

- **Turnaround time.** From the point of view of a particular process, the important criterion is how long it takes to execute that process. The interval from the time of submission of a process to the time of completion is the turnaround time. Turnaround time is the sum of the periods spent waiting in the ready queue, executing on the CPU, and doing I/O.
- **Waiting time.** The CPU-scheduling algorithm does not affect the amount of time during which a process executes or does I/O. It affects only the amount of time that a process spends waiting in the ready queue. Waiting time is the sum of the periods spent waiting in the ready queue.
- **Response time.** In an interactive system, turnaround time may not be the best criterion. Often, a process can produce some output fairly early and can continue computing new results while previous results are being output to the user. Thus, another measure is the time from the submission of a request until the first response is produced. This measure, called response time, is the time it takes to start responding, not the time it takes to output the response.

It is desirable to maximize CPU utilization and throughput and to minimize turnaround time, waiting time, and response time. In most cases, we optimize the average measure. However, under some circumstances, we prefer to optimize the minimum or maximum values rather than the average. For example, to guarantee that all users get good service, we may want to minimize the maximum response time.

Investigators have suggested that, for interactive systems (such as a PC desktop or laptop system), it is more important to minimize the variance in the response time than to minimize the average response time. A system with reasonable and predictable response time may be considered more desirable than a system that is faster on the average but is highly variable. However, little work has been done on CPU-scheduling algorithms that minimize variance.

As we discuss various CPU-scheduling algorithms in the following section, we illustrate their operation. An accurate illustration should involve many processes, each a sequence of several hundred CPU bursts and I/O bursts. For simplicity, though, we consider only one CPU burst (in milliseconds) per process in our examples. Our measure of comparison is the average waiting time. More elaborate evaluation mechanisms are discussed in Section 5.8.

5.3 Scheduling Algorithms

CPU scheduling deals with the problem of deciding which of the processes in the ready queue is to be allocated the CPU's core. There are many different CPU-scheduling algorithms. In this section, we describe several of them. Although most modern CPU architectures have multiple processing cores, we describe these scheduling algorithms in the context of only one processing core available. That is, a single CPU that has a single processing core, thus the system is

capable of only running one process at a time. In Section 5.5 we discuss CPU scheduling in the context of multiprocessor systems.


5.3.1 First-Come, First-Served Scheduling

By far the simplest CPU-scheduling algorithm is the **first-come first-serve (FCFS)** scheduling algorithm. With this scheme, the process that requests the CPU first is allocated the CPU first. The implementation of the FCFS policy is easily managed with a FIFO queue. When a process enters the ready queue, its PCB is linked onto the tail of the queue. When the CPU is free, it is allocated to the process at the head of the queue. The running process is then removed from the queue. The code for FCFS scheduling is simple to write and understand.


On the negative side, the average waiting time under the FCFS policy is often quite long. Consider the following set of processes that arrive at time 0, with the length of the CPU burst given in milliseconds:

Process	Burst Time
P_1	24
P_2	3
P_3	3

If the processes arrive in the order P_1, P_2, P_3 , and are served in FCFS order, we get the result shown in the following **Gantt chart**, which is a bar chart that illustrates a particular schedule, including the start and finish times of each of the participating processes:



The waiting time is 0 milliseconds for process P_1 , 24 milliseconds for process P_2 , and 27 milliseconds for process P_3 . Thus, the average waiting time is $(0 + 24 + 27)/3 = 17$ milliseconds. If the processes arrive in the order P_2, P_3, P_1 , however, the results will be as shown in the following Gantt chart:



The average waiting time is now $(6 + 0 + 3)/3 = 3$ milliseconds. This reduction is substantial. Thus, the average waiting time under an FCFS policy is generally not minimal and may vary substantially if the processes' CPU burst times vary greatly.

In addition, consider the performance of FCFS scheduling in a dynamic situation. Assume we have one CPU-bound process and many I/O-bound processes. As the processes flow around the system, the following scenario may result. The CPU-bound process will get and hold the CPU. During this time, all the other processes will finish their I/O and will move into the ready queue, waiting for the CPU. While the processes wait in the ready queue, the I/O

devices are idle. Eventually, the CPU-bound process finishes its CPU burst and moves to an I/O device. All the I/O-bound processes, which have short CPU bursts, execute quickly and move back to the I/O queues. At this point, the CPU sits idle. The CPU-bound process will then move back to the ready queue and be allocated the CPU. Again, all the I/O processes end up waiting in the ready queue until the CPU-bound process is done. There is a **convoy effect** as all the other processes wait for the one big process to get off the CPU. This effect results in lower CPU and device utilization than might be possible if the shorter processes were allowed to go first.

Note also that the FCFS scheduling algorithm is nonpreemptive. Once the CPU has been allocated to a process, that process keeps the CPU until it releases the CPU, either by terminating or by requesting I/O. The FCFS algorithm is thus particularly troublesome for interactive systems, where it is important that each process get a share of the CPU at regular intervals. It would be disastrous to allow one process to keep the CPU for an extended period.


5.3.2 Shortest-Job-First Scheduling

A different approach to CPU scheduling is the **shortest-job-first (SJF)** scheduling algorithm. This algorithm associates with each process the length of the process's next CPU burst. When the CPU is available, it is assigned to the process that has the smallest next CPU burst. If the next CPU bursts of two processes are the same, FCFS scheduling is used to break the tie. Note that a more appropriate term for this scheduling method would be the **shortest-next-CPU-burst** algorithm, because scheduling depends on the length of the next CPU burst of a process, rather than its total length. We use the term SJF because most people and textbooks use this term to refer to this type of scheduling.

As an example of SJF scheduling, consider the following set of processes, with the length of the CPU burst given in milliseconds:

Process	Burst Time
P_1	6
P_2	8
P_3	7
P_4	3

Using SJF scheduling, we would schedule these processes according to the following Gantt chart:



The waiting time is 3 milliseconds for process P_1 , 16 milliseconds for process P_2 , 9 milliseconds for process P_3 , and 0 milliseconds for process P_4 . Thus, the average waiting time is $(3 + 16 + 9 + 0)/4 = 7$ milliseconds. By comparison, if we were using the FCFS scheduling scheme, the average waiting time would be 10.25 milliseconds.

The SJF scheduling algorithm is provably optimal, in that it gives the minimum average waiting time for a given set of processes. Moving a short process


before a long one decreases the waiting time of the short process more than it increases the waiting time of the long process. Consequently, the average waiting time decreases.

Although the SJF algorithm is optimal, it cannot be implemented at the level of CPU scheduling, as there is no way to know the length of the next CPU burst. One approach to this problem is to try to approximate SJF scheduling. We may not know the length of the next CPU burst, but we may be able to predict its value. We expect that the next CPU burst will be similar in length to the previous ones. By computing an approximation of the length of the next CPU burst, we can pick the process with the shortest predicted CPU burst.

The next CPU burst is generally predicted as an **exponential average** of the measured lengths of previous CPU bursts. We can define the exponential average with the following formula. Let t_n be the length of the n th CPU burst, and let τ_{n+1} be our predicted value for the next CPU burst. Then, for α , $0 \leq \alpha \leq 1$, define

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$

The value of t_n contains our most recent information, while τ_n stores the past history. The parameter α controls the relative weight of recent and past history in our prediction. If $\alpha = 0$, then $\tau_{n+1} = \tau_n$, and recent history has no effect (current conditions are assumed to be transient). If $\alpha = 1$, then $\tau_{n+1} = t_n$, and only the most recent CPU burst matters (history is assumed to be old and irrelevant). More commonly, $\alpha = 1/2$, so recent history and past history are equally weighted. The initial τ_0 can be defined as a constant or as an overall system average. Figure 5.4 shows an exponential average with $\alpha = 1/2$ and $\tau_0 = 10$.



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Figure 5.4 Prediction of the length of the next CPU burst.

To understand the behavior of the exponential average, we can expand the formula for τ_{n+1} by substituting for τ_n to find

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \cdots + (1 - \alpha)^j \alpha t_{n-j} + \cdots + (1 - \alpha)^{n+1} \tau_0.$$


Typically, α is less than 1. As a result, $(1 - \alpha)$ is also less than 1, and each successive term has less weight than its predecessor.

The SJF algorithm can be either preemptive or nonpreemptive. The choice arises when a new process arrives at the ready queue while a previous process is still executing. The next CPU burst of the newly arrived process may be shorter than what is left of the currently executing process. A preemptive SJF algorithm will preempt the currently executing process, whereas a nonpreemptive SJF algorithm will allow the currently running process to finish its CPU burst. Preemptive SJF scheduling is sometimes called **shortest-remaining-time-first** scheduling.

As an example, consider the following four processes, with the length of the CPU burst given in milliseconds:

Process	Arrival Time	Burst Time
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

If the processes arrive at the ready queue at the times shown and need the indicated burst times, then the resulting preemptive SJF schedule is as depicted in the following Gantt chart:



Process P_1 is started at time 0, since it is the only process in the queue. Process P_2 arrives at time 1. The remaining time for process P_1 (7 milliseconds) is larger than the time required by process P_2 (4 milliseconds), so process P_1 is preempted, and process P_2 is scheduled. The average waiting time for this example is $[(10 - 1) + (1 - 1) + (17 - 2) + (5 - 3)]/4 = 26/4 = 6.5$ milliseconds. Nonpreemptive SJF scheduling would result in an average waiting time of 7.75 milliseconds.

5.3.3 Round-Robin Scheduling

The **round-robin (RR)** scheduling algorithm is similar to FCFS scheduling, but preemption is added to enable the system to switch between processes. A small unit of time, called a **time quantum** or **time slice**, is defined. A time quantum is generally from 10 to 100 milliseconds in length. The ready queue is treated as a circular queue. The CPU scheduler goes around the ready queue, allocating the CPU to each process for a time interval of up to 1 time quantum.

To implement RR scheduling, we again treat the ready queue as a FIFO queue of processes. New processes are added to the tail of the ready queue.

The CPU scheduler picks the first process from the ready queue, sets a timer to interrupt after 1 time quantum, and dispatches the process.

One of two things will then happen. The process may have a CPU burst of less than 1 time quantum. In this case, the process itself will release the CPU voluntarily. The scheduler will then proceed to the next process in the ready queue. If the CPU burst of the currently running process is longer than 1 time quantum, the timer will go off and will cause an interrupt to the operating system. A context switch will be executed, and the process will be put at the tail of the ready queue. The CPU scheduler will then select the next process in the ready queue.

The average waiting time under the RR policy is often long. Consider the following set of processes that arrive at time 0, with the length of the CPU burst given in milliseconds:

Process	Burst Time
P_1	24
P_2	3
P_3	3

If we use a time quantum of 4 milliseconds, then process P_1 gets the first 4 milliseconds. Since it requires another 20 milliseconds, it is preempted after the first time quantum, and the CPU is given to the next process in the queue, process P_2 . Process P_2 does not need 4 milliseconds, so it quits before its time quantum expires. The CPU is then given to the next process, process P_3 . Once each process has received 1 time quantum, the CPU is returned to process P_1 for an additional time quantum. The resulting RR schedule is as follows:

P_1	P_2	P_3	P_1	P_1	P_1	P_1	P_1	
0	4	7	10	14	18	22	26	30

Let's calculate the average waiting time for this schedule. P_1 waits for 6 milliseconds ($10 - 4$), P_2 waits for 4 milliseconds, and P_3 waits for 7 milliseconds. Thus, the average waiting time is $17/3 = 5.66$ milliseconds.

In the RR scheduling algorithm, no process is allocated the CPU for more than 1 time quantum in a row (unless it is the only runnable process). If a process's CPU burst exceeds 1 time quantum, that process is preempted and is put back in the ready queue. The RR scheduling algorithm is thus preemptive.

If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units. Each process must wait no longer than $(n - 1) \times q$ time units until its next time quantum. For example, with five processes and a time quantum of 20 milliseconds, each process will get up to 20 milliseconds every 100 milliseconds.

The performance of the RR algorithm depends heavily on the size of the time quantum. At one extreme, if the time quantum is extremely large, the RR policy is the same as the FCFS policy. In contrast, if the time quantum is extremely small (say, 1 millisecond), the RR approach can result in a large




Figure 5.5 How a smaller time quantum increases context switches.

number of context switches. Assume, for example, that we have only one process of 10 time units. If the quantum is 12 time units, the process finishes in less than 1 time quantum, with no overhead. If the quantum is 6 time units, however, the process requires 2 quanta, resulting in a context switch. If the time quantum is 1 time unit, then nine context switches will occur, slowing the execution of the process accordingly (Figure 5.5).

Thus, we want the time quantum to be large with respect to the context-switch time. If the context-switch time is approximately 10 percent of the time quantum, then about 10 percent of the CPU time will be spent in context switching. In practice, most modern systems have time quanta ranging from 10 to 100 milliseconds. The time required for a context switch is typically less than 10 microseconds; thus, the context-switch time is a small fraction of the time quantum.

Turnaround time also depends on the size of the time quantum. As we can see from Figure 5.6, the average turnaround time of a set of processes does not necessarily improve as the time-quantum size increases. In general, the average turnaround time can be improved if most processes finish their next CPU burst in a single time quantum. For example, given three processes of 10 time units each and a quantum of 1 time unit, the average turnaround time is 29. If the time quantum is 10, however, the average turnaround time drops to 20. If context-switch time is added in, the average turnaround time increases even more for a smaller time quantum, since more context switches are required.

Although the time quantum should be large compared with the context-switch time, it should not be too large. As we pointed out earlier, if the time quantum is too large, RR scheduling degenerates to an FCFS policy. A rule of thumb is that 80 percent of the CPU bursts should be shorter than the time quantum.

5.3.4 Priority Scheduling

The SJF algorithm is a special case of the general **priority-scheduling** algorithm. A priority is associated with each process, and the CPU is allocated to the




Figure 5.6 How turnaround time varies with the time quantum.

process with the highest priority. Equal-priority processes are scheduled in FCFS order. An SJF algorithm is simply a priority algorithm where the priority (p) is the inverse of the (predicted) next CPU burst. The larger the CPU burst, the lower the priority, and vice versa.

Note that we discuss scheduling in terms of *high* priority and *low* priority. Priorities are generally indicated by some fixed range of numbers, such as 0 to 7 or 0 to 4,095. However, there is no general agreement on whether 0 is the highest or lowest priority. Some systems use low numbers to represent low priority; others use low numbers for high priority. This difference can lead to confusion. In this text, we assume that low numbers represent high priority.

As an example, consider the following set of processes, assumed to have arrived at time 0 in the order $P_1, P_2, \dots, P_5$, with the length of the CPU burst given in milliseconds:

Process	Burst Time	Priority
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

Using priority scheduling, we would schedule these processes according to the following Gantt chart:



The average waiting time is 8.2 milliseconds.

Priorities can be defined either internally or externally. Internally defined priorities use some measurable quantity or quantities to compute the priority of a process. For example, time limits, memory requirements, the number of open files, and the ratio of average I/O burst to average CPU burst have been used in computing priorities. External priorities are set by criteria outside the operating system, such as the importance of the process, the type and amount of funds being paid for computer use, the department sponsoring the work, and other, often political, factors.

Priority scheduling can be either preemptive or nonpreemptive. When a process arrives at the ready queue, its priority is compared with the priority of the currently running process. A preemptive priority scheduling algorithm will preempt the CPU if the priority of the newly arrived process is higher than the priority of the currently running process. A nonpreemptive priority scheduling algorithm will simply put the new process at the head of the ready queue.


A major problem with priority scheduling algorithms is **indefinite blocking**, or **starvation**. A process that is ready to run but waiting for the CPU can be considered blocked. A priority scheduling algorithm can leave some low-priority processes waiting indefinitely. In a heavily loaded computer system, a steady stream of higher-priority processes can prevent a low-priority process from ever getting the CPU. Generally, one of two things will happen. Either the process will eventually be run (at 2 A.M. Sunday, when the system is finally lightly loaded), or the computer system will eventually crash and lose all unfinished low-priority processes. (Rumor has it that when they shut down the IBM 7094 at MIT in 1973, they found a low-priority process that had been submitted in 1967 and had not yet been run.)

A solution to the problem of indefinite blockage of low-priority processes is **aging**. Aging involves gradually increasing the priority of processes that wait in the system for a long time. For example, if priorities range from 127 (low) to 0 (high), we could periodically (say, every second) increase the priority of a waiting process by 1. Eventually, even a process with an initial priority of 127 would have the highest priority in the system and would be executed. In fact, it would take a little over 2 minutes for a priority-127 process to age to a priority-0 process.

Another option is to combine round-robin and priority scheduling in such a way that the system executes the highest-priority process and runs processes with the same priority using round-robin scheduling. Let's illustrate with an example using the following set of processes, with the burst time in milliseconds:

Process	Burst Time	Priority
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3

Using priority scheduling with round-robin for processes with equal priority, we would schedule these processes according to the following Gantt chart using a time quantum of 2 milliseconds:



In this example, process P_4 has the highest priority, so it will run to completion. Processes P_2 and P_3 have the next-highest priority, and they will execute in a round-robin fashion. Notice that when process P_2 finishes at time 16, process P_3 is the highest-priority process, so it will run until it completes execution. Now, only processes P_1 and P_5 remain, and as they have equal priority, they will execute in round-robin order until they complete.

5.3.5 Multilevel Queue Scheduling

With both priority and round-robin scheduling, all processes may be placed in a single queue, and the scheduler then selects the process with the highest priority to run. Depending on how the queues are managed, an $O(n)$ search may be necessary to determine the highest-priority process. In practice, it is often easier to have separate queues for each distinct priority, and priority scheduling simply schedules the process in the highest-priority queue. This is illustrated in Figure 5.7. This approach—known as **multilevel queue**—also works well when priority scheduling is combined with round-robin: if there are multiple processes in the highest-priority queue, they are executed in round-robin order. In the most generalized form of this approach, a priority is assigned statically to each process, and a process remains in the same queue for the duration of its runtime.




Figure 5.7 Separate queues for each priority.




Figure 5.8 Multilevel queue scheduling.

A multilevel queue scheduling algorithm can also be used to partition processes into several separate queues based on the process type (Figure 5.8). For example, a common division is made between **foreground** (interactive) processes and **background** (batch) processes. These two types of processes have different response-time requirements and so may have different scheduling needs. In addition, foreground processes may have priority (externally defined) over background processes. Separate queues might be used for foreground and background processes, and each queue might have its own scheduling algorithm. The foreground queue might be scheduled by an RR algorithm, for example, while the background queue is scheduled by an FCFS algorithm.

In addition, there must be scheduling among the queues, which is commonly implemented as fixed-priority preemptive scheduling. For example, the real-time queue may have absolute priority over the interactive queue.

Let's look at an example of a multilevel queue scheduling algorithm with four queues, listed below in order of priority:

1. Real-time processes
2. System processes
3. Interactive processes
4. Batch processes

Each queue has absolute priority over lower-priority queues. No process in the batch queue, for example, could run unless the queues for real-time processes, system processes, and interactive processes were all empty. If an interactive process entered the ready queue while a batch process was running, the batch process would be preempted.

Another possibility is to time-slice among the queues. Here, each queue gets a certain portion of the CPU time, which it can then schedule among its various processes. For instance, in the foreground–background queue example, the foreground queue can be given 80 percent of the CPU time for RR scheduling

among its processes, while the background queue receives 20 percent of the CPU to give to its processes on an FCFS basis.

5.3.6 Multilevel Feedback Queue Scheduling

Normally, when the multilevel queue scheduling algorithm is used, processes are permanently assigned to a queue when they enter the system. If there are separate queues for foreground and background processes, for example, processes do not move from one queue to the other, since processes do not change their foreground or background nature. This setup has the advantage of low scheduling overhead, but it is inflexible.

The **multilevel feedback queue** scheduling algorithm, in contrast, allows a process to move between queues. The idea is to separate processes according to the characteristics of their CPU bursts. If a process uses too much CPU time, it will be moved to a lower-priority queue. This scheme leaves I/O-bound and interactive processes—which are typically characterized by short CPU bursts—in the higher-priority queues. In addition, a process that waits too long in a lower-priority queue may be moved to a higher-priority queue. This form of aging prevents starvation.

For example, consider a multilevel feedback queue scheduler with three queues, numbered from 0 to 2 (Figure 5.9). The scheduler first executes all processes in queue 0. Only when queue 0 is empty will it execute processes in queue 1. Similarly, processes in queue 2 will be executed only if queues 0 and 1 are empty. A process that arrives for queue 1 will preempt a process in queue 2. A process in queue 1 will in turn be preempted by a process arriving for queue 0.

An entering process is put in queue 0. A process in queue 0 is given a time quantum of 8 milliseconds. If it does not finish within this time, it is moved to the tail of queue 1. If queue 0 is empty, the process at the head of queue 1 is given a quantum of 16 milliseconds. If it does not complete, it is preempted and is put into queue 2. Processes in queue 2 are run on an FCFS basis but are run only when queues 0 and 1 are empty. To prevent starvation, a process that waits too long in a lower-priority queue may gradually be moved to a higher-priority queue.




Figure 5.9 Multilevel feedback queues.

This scheduling algorithm gives highest priority to any process with a CPU burst of 8 milliseconds or less. Such a process will quickly get the CPU, finish its CPU burst, and go off to its next I/O burst. Processes that need more than 8 but less than 24 milliseconds are also served quickly, although with lower priority than shorter processes. Long processes automatically sink to queue 2 and are served in FCFS order with any CPU cycles left over from queues 0 and 1.

In general, a multilevel feedback queue scheduler is defined by the following parameters:

- The number of queues
- The scheduling algorithm for each queue
- The method used to determine when to upgrade a process to a higher-priority queue
- The method used to determine when to demote a process to a lower-priority queue
- The method used to determine which queue a process will enter when that process needs service

The definition of a multilevel feedback queue scheduler makes it the most general CPU-scheduling algorithm. It can be configured to match a specific system under design. Unfortunately, it is also the most complex algorithm, since defining the best scheduler requires some means by which to select values for all the parameters.

5.4 Thread Scheduling

In Chapter 4, we introduced threads to the process model, distinguishing between *user-level* and *kernel-level* threads. On most modern operating systems it is kernel-level threads—not processes—that are being scheduled by the operating system. User-level threads are managed by a thread library, and the kernel is unaware of them. To run on a CPU, user-level threads must ultimately be mapped to an associated kernel-level thread, although this mapping may be indirect and may use a lightweight process (LWP). In this section, we explore scheduling issues involving user-level and kernel-level threads and offer specific examples of scheduling for Pthreads.

5.4.1 Contention Scope

One distinction between user-level and kernel-level threads lies in how they are scheduled. On systems implementing the many-to-one (Section 4.3.1) and many-to-many (Section 4.3.3) models, the thread library schedules user-level threads to run on an available LWP. This scheme is known as **process-contention scope (PCS)**, since competition for the CPU takes place among threads belonging to the same process. (When we say the thread library *schedules* user threads onto available LWPs, we do not mean that the threads are actually *running* on a CPU as that further requires the operating system to schedule the LWP's kernel thread onto a physical CPU core.) To decide which

Student Name: سمير محمد عباس حيدر	الاسم الكامل باللغة العربية	
نظم تشغيل	أسم المادة	
Course Instructor د.لوي	اسم التدريسي	
Study: صباحي	صباحي او مسائي	
Stage: ثالثة	المرحلة	
Submission Date:	تاريخ التسليم	

Comprehensive Report: CPU Scheduling Algorithms in Operating Systems

1. Introduction to CPU Scheduling

The Central Processing Unit (CPU) is one of the most valuable resources in any computer system. Within operating systems, scheduling is the fundamental decision-making process that determines which process gets access to the CPU next and for how long. The Short-Term Scheduler, also known as the CPU scheduler, makes these critical decisions dozens or even hundreds of times per second.

During execution, a process typically alternates between CPU bursts (executing instructions) and I/O bursts (waiting for input/output operations). Scheduling occurs at every transition between these states, directly shaping overall system performance, responsiveness, and fairness.

2. Core Scheduling Objectives

Every scheduling algorithm attempts to balance multiple competing goals, though no single algorithm can optimize all of them simultaneously. The primary objectives include:

- **Performance:** The system must maximize throughput by completing as many processes per unit time as possible. It also aims to maximize overall CPU utilization.
- **Responsiveness:** It is crucial to minimize wait times (the duration processes sit in the ready queue before getting CPU time) and response times. Low response time is especially critical for interactive systems where users expect quick feedback.
- **Fairness:** The scheduler must ensure that no process is starved indefinitely, meaning every process should eventually gain CPU access.

3. Primary Scheduling Algorithms

3.1. First-Come, First-Served (FCFS)

FCFS is the simplest scheduling strategy available.

- Processes are executed strictly in their arrival order, functioning like a standard queue.

- It is a non-preemptive algorithm; once a process begins running, it continues until completion.
- **Strengths:** It is highly simple to implement and guarantees no starvation.
- **Weaknesses:** It suffers from the "convoy effect," a scenario where short processes are forced to wait behind long ones, drastically inflating the average wait time.

3.2. Shortest Job First (SJF)

The SJF algorithm selects the process with the shortest predicted CPU burst to run next.

- It exists in two primary forms: non-preemptive (runs to completion) and preemptive (known as Shortest Remaining Time First, or SRTF, where a newly arrived shorter process can interrupt the currently running one).
- Mathematically, SJF is provably optimal for minimizing average waiting time.
- However, it requires knowing the burst lengths in advance, which is typically estimated using exponential averaging.
- A major drawback is the risk of starvation, as long jobs may never execute if shorter jobs continuously arrive.

3.3. Priority Scheduling

In Priority Scheduling, each process is assigned a specific priority value.

- The CPU is always allocated to the ready process with the highest priority, and any ties are typically broken using FCFS.
- Priorities can be "Internal" (derived from process properties like memory needs or I/O ratios) or "External" (set by the user or system policies).
- To solve the inherent starvation problem associated with priority scheduling, modern systems use "Aging". Aging gradually increases the priority of waiting processes over time to ensure they eventually run.

3.4. Round Robin (RR)

Round Robin is fundamentally designed for fairness and time-sharing environments.

- Each process receives a fixed time quantum (typically between 10 and 100 ms).
- After a process's slice expires, it is safely preempted and moved to the back of the queue.
- This ensures all processes get equal CPU turns, completely eliminating starvation.
- The size of the quantum is highly critical: if it is too large, the system behaves like FCFS, and if

it is too small, the system suffers from high context-switch overhead. RR provides excellent response times for interactive and multi-user systems.

4. Algorithm Comparison Matrix

Algorithm	Preemptive?	Avg. Wait Time	Starvation Risk?	Best Suited For
FCFS	No	High	No	Batch systems
SJF	Optional	Lowest	Yes	Batch / offline processes
Priority	Optional	Moderate	Yes	Real-time systems
Round Robin	Yes	Moderate	No	Interactive / time-sharing

5. Conclusion and Best Practices

When designing or configuring an operating system, there is no universal "best" scheduler. The optimal choice heavily depends on the system's workload and operational goals. While SJF is mathematically optimal for reducing waiting times, Round Robin is superior regarding fairness.

To select the right algorithm, system architects must:

1. **Identify Workload Type:** Determine whether the system handles batch, interactive, real-time, or mixed processes.
2. **Define Performance Goals:** Decide whether to prioritize throughput, response time, or fairness.
3. **Evaluate Trade-offs:** Consider factors like starvation risk, overhead limits, and implementation complexity.

Because single algorithms have distinct weaknesses, modern operating system kernels typically employ hybrid solutions—such as multilevel feedback queues—to effectively balance these trade-offs and yield the best overall performance.

Sources & References:

SYSTEM CALLS AND THEIR SECURITY

A Comprehensive Technical Report

Abdullah Qusay Mejbel • Bashar Haider Moneum • Mohammed Adnan Abbas

1. Introduction to System Calls

System calls function as the essential intermediary between user applications and the computer's hardware. Because direct hardware access is strictly prohibited for security reasons, system calls provide a controlled entry point for applications to request kernel services.

They serve as the "API of the Kernel," ensuring hardware abstraction and the uniform enforcement of security policies across all running processes.

2. The Mechanics of the User-Kernel Boundary

Operating systems operate using a dual-mode system enforced by hardware, separating privilege levels to protect system integrity:

- **User Mode (Ring 3):** Applications run with a limited instruction set and isolated memory to prevent them from crashing the system.
- **Kernel Mode (Ring 0):** The kernel has unrestricted access to the CPU and hardware memory.

2.1 The Trap Mechanism

When an application requires a kernel service, it triggers a "trap" or software interrupt using specific instructions — such as `syscall` for x64 architecture. This causes the processor to switch from User Mode to Kernel Mode, transferring control to a predefined handler. The kernel then uses a System Call Table to look up the unique ID of the request and locate the corresponding service routine address.




Figure 1.13 Transition from user to kernel mode.

Figure 1.13 — Transition from user to kernel mode.

Step	Actor	Description
1	Application	Invokes syscall instruction with syscall ID in register
2	CPU	Switches from Ring 3 (User) to Ring 0 (Kernel)
3	Kernel	Looks up syscall ID in the System Call Table
4	Kernel	Executes the corresponding service routine
5	CPU	Returns to User Mode with result

3. Categories of System Calls

System calls are generally categorized by the type of service they provide to user-space applications:

Category	Key Calls	Purpose
Process Control	fork(), exit(), wait()	Create, terminate, and manage processes
File Management	open(), read(), write(), close()	Perform file I/O operations
Device Management	ioctl(), read(), write()	Manage character or block devices
Information Maintenance	getpid(), time(), uname()	Retrieve system and process metadata
Communication (IPC)	pipe(), socket(), shmget()	Facilitate data flow between processes
Protection & Security	chmod(), chown(), setuid()	Control identities and permissions

1. Introduction

An Operating System (OS) is the software that runs everything in a computer. It controls hardware, runs applications, and manages files and memory. Because of this central role, the OS is the most important layer when it comes to security.

If an attacker manages to reach the OS level, every application and every piece of data becomes vulnerable. This is why OS security is not just a technical topic — it is a critical responsibility for every developer, student, and user.

This report covers the main concepts of OS security in a simple and clear way. It explains the difference between security and protection, describes the main threats that exist today, and shows how modern operating systems defend against them.

2. Security vs. Protection

People often use the words 'security' and 'protection' as if they mean the same thing, but in computer science they have different meanings that are important to understand.

Protection refers to the rules inside the OS that control which users or programs can access which resources. Think of it as the internal rules of a building — who is allowed into which room.

Security refers to the broader effort to defend the system against outside threats like hackers, viruses, or physical attacks. Think of it as the fence, the alarm system, and the security cameras around the building.

Feature	Protection	Security
Focus	Inside the OS — between users and programs	Outside the OS — against hackers and malware
Goal	Prevent unauthorized internal access	Defend against intentional external attacks
Example	User A cannot delete User B's files	A firewall blocks a remote attacker
If it fails	Accidental data loss or internal misuse	Data breach, ransomware, system takeover

Important point: Protection can still be working perfectly while Security has already failed. If an attacker steals a user's password and logs in normally, the protection rules still work — but the system is not secure at all.

3. The CIA Triad and Security Principles

The CIA Triad is the most important framework in information security. Every security decision should be measured against these three properties:

- **Confidentiality:** Only authorized people can see sensitive information. This is achieved through encryption (like AES-256) and strict access controls.
- **Integrity:** Data cannot be changed by unauthorized people. Cryptographic hashes and digital signatures verify that files have not been tampered with.
- **Availability:** The system must be accessible when needed. Defending against Denial of Service (DoS) attacks is the main concern here.

In addition to the CIA Triad, secure OS design follows these principles:

- **Least Privilege:** Every user and program should have only the minimum access they need. This limits damage if something goes wrong.
 - **Defense in Depth:** Use multiple layers of security. If one layer fails, the next one still protects the system.
 - **Fail-Safe Defaults:** Access should be denied by default. Permissions must be granted explicitly, never assumed.
 - **Separation of Privilege:** Critical actions should require more than one condition. Multi-Factor Authentication (MFA) is a direct example of this principle.
-

4. Types of Threats and Attacks

4.1 Malware Types

Malware is any software designed to cause harm. There are five main types:

- **Virus:** Attaches to a program and spreads when a user runs that program. Requires human action to activate.
- **Worm:** Spreads automatically across networks without any user action by exploiting OS vulnerabilities.
- **Trojan:** Looks like a useful program but hides malicious code. Unlike worms, it does not self-replicate.
- **Rootkit:** Hides deep inside the OS, often in the kernel, making it extremely difficult to detect. Gives attackers persistent hidden access.
- **Ransomware:** Encrypts all of the victim's files and demands payment to restore them. WannaCry is the most famous example.

4.2 Technical Exploits

- **Buffer Overflow:** When a program writes more data into memory than it can hold, the extra data can overwrite other parts of memory and let an attacker run their own code.
- **Privilege Escalation:** An attacker with limited access exploits a bug to gain administrator-level control over the system.
- **Side-Channel Attacks:** Attacks like Spectre and Meltdown exploit the physical behavior of the CPU to read private data, bypassing OS protection entirely.

4.3 Attack Types and What They Target

Attack Type	What It Does	CIA Property Violated
Interception	Secretly reads data — e.g. wiretapping	Confidentiality
Modification	Changes data without permission	Integrity
Interruption	Shuts down a service — DoS attack	Availability
Fabrication	Inserts fake data into the system	Integrity and Authenticity

What is a Thread

-A thread is the smallest unit of execution within a process.





Thread Scheduling

-Thread scheduling is how
An operating system decides
which thread runs next.



we defined **deadlock** as a situation in which every **process** in a set of **processes** is waiting for an event that can be caused only by another **process** in the set.

2. Deadlock Conditions

A **deadlock** situation can arise if the following four conditions hold simultaneously in a system:

- **Mutual exclusion.** At least one **resource** must be held in a **non-sharable mode**; that is, only one **thread** at a time can use the **resource**. If another **thread** requests that **resource**, the requesting **thread** must be delayed until the **resource** has been released.
- **Hold and wait.** A **thread** must be holding at least one **resource** and waiting to acquire additional **resources** that are currently being held by other **threads**.
- **No preemption.** **Resources** cannot be **preempted**; that is, a **resource** can be released only voluntarily by the **thread** holding it, after that **thread** has completed its task.
- **Circular wait.** A set $\{T_0, T_1, \dots, T_n\}$ of **waiting threads** must exist such that T_0 is waiting for a **resource** held by T_1 , T_1 is waiting for a **resource** held by T_2 , ..., T_{n-1} is waiting for a **resource** held by T_n , and T_n is waiting for a **resource** held by T_0 .

We emphasize that all four conditions must hold for a **deadlock** to occur. The **circular-wait** condition implies the **hold-and-wait** condition, so the four conditions are not completely independent. Note, however, that it is useful to consider each condition separately.

2. What is a Thread?

A thread is the smallest schedulable unit of CPU execution within an operating system. It represents a single sequential flow of control inside a process and is often described as a lightweight process because it shares most of its resources with other threads in the same process, yet it can be independently scheduled by the CPU.

In a multi-threaded system, a single process can contain multiple threads that execute concurrently — either running different parts of a program or cooperating on the same task. This concurrency improves performance by allowing better CPU utilization and reducing idle time, especially on multi-core systems where threads can run simultaneously on separate cores.

7. Synchronization & Concurrency Control

7.1 The Problem: Race Conditions

A race condition occurs when two or more threads access shared data simultaneously and the final result depends on the unpredictable order of their execution. Even simple operations like incrementing a counter (`counter++`) can be non-atomic at the machine code level, involving three steps: read, increment, write — any of which can be interrupted.

Example: Two threads each execute `counter++` concurrently on an initial value of 5. Due to interleaving, the result may be 6 instead of the expected 7 — this is a race condition caused by unsynchronized shared access.

7.2 Deadlocks

A deadlock is a state in which two or more threads are permanently blocked, each waiting for a resource held by another, creating a circular dependency with no way to resolve it.

Four necessary conditions must all hold simultaneously for a deadlock to occur (Coffman Conditions):

- Mutual Exclusion: At least one resource can only be held by one thread at a time.
- Hold and Wait: A thread holding at least one resource is waiting to acquire additional resources held by others.
- No Preemption: Resources cannot be forcibly taken from a thread — they must be released voluntarily.
- Circular Wait: A set of threads $T_1, T_2, \dots, T_n$ exists such that T_1 is waiting for a resource held by T_2 , T_2 is waiting for T_3 , and so on, with T_n waiting for T_1 .

7.3 Starvation

Starvation occurs when a thread is indefinitely denied access to the resources it needs, typically because higher-priority threads or unfair scheduling policies continuously preempt it. Unlike deadlock, the thread is not blocked — it is simply never selected. The solution is aging: gradually elevating the priority of long-waiting threads.

8. Advantages of Multithreading

Multithreading provides several key performance and design advantages in modern operating systems. One of the most important benefits is responsiveness, especially in graphical user interface (GUI) applications, where one thread can handle user interactions while another performs intensive computations, preventing the system from freezing or becoming unresponsive.

It also improves resource sharing because threads within the same process operate in a shared memory space, eliminating the need for complex inter-process communication (IPC) mechanisms. This shared environment reduces overhead and simplifies data exchange between execution units.

From a performance perspective, multithreading is economically efficient since context switching between threads is significantly faster than between processes, as there is no need to switch memory address spaces. Additionally, it enables better utilization of multi-core architectures by allowing true parallelism, where multiple threads can execute simultaneously on different CPU cores, increasing overall system throughput.
